

Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 1

1. Write programs to parse text files, CSV, HTML, XML and JSON documents and extract relevant data. After retrieving data check any anomalies in the data, missing values etc.

```
import pandas as pd
import numpy as np
import os

# =====
# 0. SETUP: CREATE DUMMY FILES FOR TESTING
# =====
# This generates sample files in Colab so you can test the code immediately

csv_data = """id,name,age,salary
1,Alice,28,50000
2,Bob,35,
3,Charlie,,70000
1,Alice,28,50000""" # Intentional missing values and duplicate Alice

with open('sample.csv', 'w') as f:
    f.write(csv_data)

json_data = """[
    {"id": 1, "name": "Alice", "age": 28, "salary": 50000},
    {"id": 2, "name": "Bob", "age": 35, "salary": null},
    {"id": 3, "name": "Charlie", "age": null, "salary": 70000}
]"""
with open('sample.json', 'w') as f:
    f.write(json_data)

html_data = """
<table>
  <tr><th>id</th><th>name</th><th>age</th><th>salary</th></tr>
  <tr><td>1</td><td>Alice</td><td>28</td><td>50000</td></tr>
  <tr><td>2</td><td>Bob</td><td>35</td><td></td></tr>
</table>
"""
with open('sample.html', 'w') as f:
    f.write(html_data)

xml_data = """<?xml version="1.0" ?>
<data>
  <record><id>1</id><name>Alice</name><age>28</age><salary>50000</salary></record>
</data>
</xml>"""
```

```

    <record><id>2</id><name>Bob</name><age>35</age><salary></salary></record>
</data>
"""
with open('sample.xml', 'w') as f:
    f.write(xml_data)

# =====
# 1. PARSING FUNCTIONS
# =====

def parse_csv(file_path):
    print(f"\n--- Parsing CSV: {file_path} ---")
    return pd.read_csv(file_path)

def parse_json(file_path):
    print(f"\n--- Parsing JSON: {file_path} ---")
    return pd.read_json(file_path)

def parse_html(file_path):
    print(f"\n--- Parsing HTML: {file_path} ---")
    tables = pd.read_html(file_path)
    return tables[0] # Returns the first table found in the HTML

def parse_xml(file_path):
    print(f"\n--- Parsing XML: {file_path} ---")
    return pd.read_xml(file_path)

def parse_text(file_path, delimiter='\t'):
    print(f"\n--- Parsing Text File: {file_path} ---")
    return pd.read_csv(file_path, delimiter=delimiter)

# =====
# 2. DATA QUALITY & ANOMALY DETECTION
# =====

def check_data_anomalies(df, dataset_name="Dataset"):
    """
    Analyzes a DataFrame for missing values, duplicates, and basic anomalies.
    """
    print(f"\n[{dataset_name}] Data Quality Report:")
    print("-" * 50)

    if df is None or df.empty:
        print("Dataframe is empty or invalid.")
        return

    # 1. Check for Missing Values (Nulls/NaNs)
    missing_values = df.isnull().sum()
    total_missing = missing_values.sum()
    if total_missing > 0:
        print(f"⚠ Missing Values Found:\n{missing_values[missing_values > 0]}\n")
    else:
        print(f"✅ No missing values found.\n")

```

```

# 2. Check for Duplicate Rows
duplicates = df.duplicated().sum()
if duplicates > 0:
    print(f"⚠ Found {duplicates} duplicate row(s).\n")
else:
    print("✅ No duplicate rows found.\n")

# 3. Show Data Types
print("Data Types:")
print(df.dtypes)
print("\n" + "=" * 50)

# =====
# 3. MAIN EXECUTION
# =====
if __name__ == "__main__":

    # 1. Test CSV
    df_csv = parse_csv('sample.csv')
    check_data_anomalies(df_csv, "CSV Data")

    # 2. Test JSON
    df_json = parse_json('sample.json')
    check_data_anomalies(df_json, "JSON Data")

    # 3. Test HTML
    df_html = parse_html('sample.html')
    check_data_anomalies(df_html, "HTML Data")

    # 4. Test XML
    df_xml = parse_xml('sample.xml')
    check_data_anomalies(df_xml, "XML Data")

```

--- Parsing CSV: sample.csv ---

[CSV Data] Data Quality Report:

⚠ Missing Values Found:

```

age      1
salary   1
dtype: int64

```

⚠ Found 1 duplicate row(s).

Data Types:

```

id      int64
name    object
age     float64
salary  float64
dtype: object

```

=====

--- Parsing JSON: sample.json ---

[JSON Data] Data Quality Report:

⚠ Missing Values Found:

age 1
salary 1
dtype: int64

✅ No duplicate rows found.

Data Types:

id int64
name object
age float64
salary float64
dtype: object

=====

--- Parsing HTML: sample.html ---

[HTML Data] Data Quality Report:

⚠ Missing Values Found:

salary 1
dtype: int64

✅ No duplicate rows found.

Data Types:

id int64
name object
age int64
salary float64
dtype: object

=====

--- Parsing XML: sample.xml ---

[XML Data] Data Quality Report:

⚠ Missing Values Found:

salary 1
dtype: int64

✅ No duplicate rows found.

Data Types:

id int64
name object
age int64
salary float64

dtype: object

=====

2. Write programs for reading and writing binary files

```
import pandas as pd
import pickle
import os

# =====
# 1. RAW BINARY READ/WRITE FUNCTIONS
# =====

def write_raw_binary(file_path, byte_data):
    """Writes raw byte data to a binary file."""
    print(f"Writing raw binary data to '{file_path}'...")
    # 'wb' stands for Write Binary
    with open(file_path, 'wb') as file:
        file.write(byte_data)
    print("✅ Write successful.\n")

def read_raw_binary(file_path):
    """Reads raw byte data from a binary file."""
    print(f"Reading raw binary data from '{file_path}'...")
    # 'rb' stands for Read Binary
    with open(file_path, 'rb') as file:
        data = file.read()
    return data

# =====
# 2. PICKLE BINARY (OBJECT SERIALIZATION)
# =====

def write_pickle_binary(file_path, data_object):
    """Serializes a Python object (like a DataFrame) to a binary file."""
    print(f"Writing DataFrame to binary pickle file '{file_path}'...")
    with open(file_path, 'wb') as file:
        pickle.dump(data_object, file)
    print("✅ Pickling successful.\n")

def read_pickle_binary(file_path):
    """Deserializes a Python object from a binary pickle file."""
    print(f"Reading DataFrame from binary pickle file '{file_path}'...")
    with open(file_path, 'rb') as file:
        data_object = pickle.load(file)
    return data_object

# =====
# 3. MAIN EXECUTION (Colab Ready)
# =====

if __name__ == "__main__":
```

```

print("=== PART 1: Raw Binary Operations ===")
raw_file = 'raw_data.bin'

# Prefixing a string with 'b' converts it to raw bytes
my_binary_data = b"Data Engineering is awesome! \x00\x01\x02\x03"

# Write and Read
write_raw_binary(raw_file, my_binary_data)
retrieved_raw_data = read_raw_binary(raw_file)

print(f"Retrieved Bytes: {retrieved_raw_data}")
print("-" * 50)

print("\n=== PART 2: DataFrame Pickling Operations ===")
pickle_file = 'dataframe_data.pkl'

# Create a sample pandas DataFrame
df = pd.DataFrame({
    'emp_id': [101, 102, 103],
    'name': ['Alice', 'Bob', 'Charlie'],
    'department': ['Data', 'Engineering', 'Analytics']
})

# Write and Read
write_pickle_binary(pickle_file, df)
retrieved_df = read_pickle_binary(pickle_file)

print("Retrieved DataFrame from Binary:")
print(retrieved_df)
print("-" * 50)

```

=== PART 1: Raw Binary Operations ===

Writing raw binary data to 'raw_data.bin'...

✓ Write successful.

Reading raw binary data from 'raw_data.bin'...

Retrieved Bytes: b'Data Engineering is awesome! \x00\x01\x02\x03'

=== PART 2: DataFrame Pickling Operations ===

Writing DataFrame to binary pickle file 'dataframe_data.pkl'...

✓ Pickling successful.

Reading DataFrame from binary pickle file 'dataframe_data.pkl'...

Retrieved DataFrame from Binary:

	emp_id	name	department
0	101	Alice	Data
1	102	Bob	Engineering
2	103	Charlie	Analytics

3. Write programs for searching, splitting, and replacing strings based on pattern matching using regular expressions

```
import re

text_data = "Contact support at support@example.com or sales at sales123@example.org. Call 555-1234."

# 1. Searching (Extracting Emails)
# Pattern looks for characters, an @ symbol, characters, a dot, and a domain suffix
email_pattern = r'[a-zA-Z0-9_+-.]+@[a-zA-Z0-9-]+\.[a-zA-Z0-9-.]+'
emails_found = re.findall(email_pattern, text_data)
print("Searched Emails:", emails_found)
# Output: ['support@example.com', 'sales123@example.org']

# 2. Splitting (Split string by spaces OR punctuation)
# Pattern splits on any whitespace or standard punctuation marks
split_pattern = r'[\s,.]+'
words = re.split(split_pattern, text_data)
# Filter out empty strings from the result
words = [w for w in words if w]
print("Split Words:", words)

# 3. Replacing (Anonymize phone numbers)
# Pattern looks for a standard 3-4 digit phone format and replaces it
phone_pattern = r'\d{3}-\d{4}'
anonymized_text = re.sub(phone_pattern, 'XXX-XXXX', text_data)
print("Replaced Text:", anonymized_text)
# Output: "...Call XXX-XXXX."
```

```
Searched Emails: ['support@example.com', 'sales123@example.org.']
Split Words: ['Contact', 'support', 'at', 'support@example', 'com', 'or', 'sales', 'at', 'sales123@example', 'org', 'Call', '555-1234']
Replaced Text: Contact support at support@example.com or sales at sales123@example.org. Call XXX-XXXX.
```

4. Design a relational database for a small application and populate the database. Using SQL do the CRUD (create, read, update and delete) operation

```
import sqlite3

# Connect to database (creates the file if it doesn't exist)
conn = sqlite3.connect('company.db')
cursor = conn.cursor()

# DESIGN: Create Table
cursor.execute('''
    CREATE TABLE IF NOT EXISTS employees (
        id INTEGER PRIMARY KEY AUTOINCREMENT,
        name TEXT NOT NULL,
        department TEXT NOT NULL,
```

```
        salary REAL
    )
'''
conn.commit()

# CREATE: Insert Data
employees_data = [
    ('Alice Smith', 'Engineering', 85000.00),
    ('Bob Jones', 'Marketing', 62000.00),
    ('Charlie Brown', 'Sales', 55000.00)
]
cursor.executemany('''
    INSERT INTO employees (name, department, salary)
    VALUES (?, ?, ?)
''', employees_data)
conn.commit()
print("Data Inserted.")

# READ: Retrieve Data
cursor.execute("SELECT * FROM employees WHERE salary > 60000")
high_earners = cursor.fetchall()
print("\nHigh Earners (Read):")
for emp in high_earners:
    print(emp)

# UPDATE: Modify Existing Data
cursor.execute('''
    UPDATE employees
    SET salary = 58000.00
    WHERE name = 'Charlie Brown'
''')
conn.commit()
print("\nRecord Updated.")

# DELETE: Remove Data
cursor.execute('''
    DELETE FROM employees
    WHERE department = 'Marketing'
''')
conn.commit()
print("Record Deleted.")

# Verify Final State
cursor.execute("SELECT * FROM employees")
print("\nFinal Database State:")
for row in cursor.fetchall():
    print(row)

# Close the connection
conn.close()
```

Data Inserted.

High Earners (Read):

```
(1, 'Alice Smith', 'Engineering', 85000.0)
(4, 'Alice Smith', 'Engineering', 85000.0)
(7, 'Alice Smith', 'Engineering', 85000.0)
(10, 'Alice Smith', 'Engineering', 85000.0)
(11, 'Bob Jones', 'Marketing', 62000.0)
```

Record Updated.

Record Deleted.

Final Database State:

```
(1, 'Alice Smith', 'Engineering', 85000.0)
(3, 'Charlie Brown', 'Sales', 58000.0)
(4, 'Alice Smith', 'Engineering', 85000.0)
(6, 'Charlie Brown', 'Sales', 58000.0)
(7, 'Alice Smith', 'Engineering', 85000.0)
(9, 'Charlie Brown', 'Sales', 58000.0)
(10, 'Alice Smith', 'Engineering', 85000.0)
(12, 'Charlie Brown', 'Sales', 58000.0)
```

Name : Rushikesh Anil Khairnar

Student ID: 5115301

Roll NO: 21

Subject: Data Engineering


Practical No. – 2

Aim:

To perform the ETL process in Power BI by extracting, transforming, and loading data to prepare it for analysis.

Step 1- Import Excel Data into Power BI

The Excel workbook containing the Products, Orders, and Order_Details tables was imported into Power BI for further data transformation and analysis.



Step 2- Open Power Query Editor

Power Query Editor was opened to perform the required data transformation operations on the imported tables.

Power Query Editor

Home Transform Add column View Help

Close & Apply, Get data, Enter data, Manage parameters, Refresh, Properties, Advanced editor, Export query results, Choose columns, Remove columns, Keep rows, Remove rows, Filter rows, Transform, Combine, Export template

Queries [2]: Orders, Products

Table.ReorderColumns(#"Added custom", {"OrderID", "CustomerID", "EmployeeID", "OrderDate", "RequiredDate", "ShippedDate", "ShipVia", "ShipDate"})

	OrderID	CustomerID	EmployeeID	OrderDate	RequiredDate	ShippedDate	ShipVia	ShipDate
1	10248	VINET	5	7/4/1996 12:00:00 AM	8/1/1996 12:00:00 AM	7/16/1996 12:00:00 AM	3	
2	10248	VINET	5	7/4/1996 12:00:00 AM	8/1/1996 12:00:00 AM	7/16/1996 12:00:00 AM	3	
3	10248	VINET	5	7/4/1996 12:00:00 AM	8/1/1996 12:00:00 AM	7/16/1996 12:00:00 AM	3	
4	10249	TOMSP	6	7/5/1996 12:00:00 AM	8/16/1996 12:00:00 AM	7/10/1996 12:00:00 AM	1	
5	10249	TOMSP	6	7/5/1996 12:00:00 AM	8/16/1996 12:00:00 AM	7/10/1996 12:00:00 AM	1	
6	10250	HANAR	4	7/8/1996 12:00:00 AM	8/5/1996 12:00:00 AM	7/12/1996 12:00:00 AM	2	
7	10250	HANAR	4	7/8/1996 12:00:00 AM	8/5/1996 12:00:00 AM	7/12/1996 12:00:00 AM	2	
8	10250	HANAR	4	7/8/1996 12:00:00 AM	8/5/1996 12:00:00 AM	7/12/1996 12:00:00 AM	2	
9	10251	VICTE	3	7/8/1996 12:00:00 AM	8/5/1996 12:00:00 AM	7/15/1996 12:00:00 AM	1	
10	10251	VICTE	3	7/8/1996 12:00:00 AM	8/5/1996 12:00:00 AM	7/15/1996 12:00:00 AM	1	
11	10251	VICTE	3	7/8/1996 12:00:00 AM	8/5/1996 12:00:00 AM	7/15/1996 12:00:00 AM	1	
12	10252	SUPRD	4	7/9/1996 12:00:00 AM	8/6/1996 12:00:00 AM	7/11/1996 12:00:00 AM	2	
13	10252	SUPRD	4	7/9/1996 12:00:00 AM	8/6/1996 12:00:00 AM	7/11/1996 12:00:00 AM	2	
14	10252	SUPRD	4	7/9/1996 12:00:00 AM	8/6/1996 12:00:00 AM	7/11/1996 12:00:00 AM	2	
15	10253	HANAR	3	7/10/1996 12:00:00 AM	7/24/1996 12:00:00 AM	7/16/1996 12:00:00 AM	2	
16	10253	HANAR	3	7/10/1996 12:00:00 AM	7/24/1996 12:00:00 AM	7/16/1996 12:00:00 AM	2	
17	10253	HANAR	3	7/10/1996 12:00:00 AM	7/24/1996 12:00:00 AM	7/16/1996 12:00:00 AM	2	
18	10254	CHOPS	5	7/11/1996 12:00:00 AM	8/8/1996 12:00:00 AM	7/23/1996 12:00:00 AM	2	
19	10254	CHOPS	5	7/11/1996 12:00:00 AM	8/8/1996 12:00:00 AM	7/23/1996 12:00:00 AM	2	
20	10254	CHOPS	5	7/11/1996 12:00:00 AM	8/8/1996 12:00:00 AM	7/23/1996 12:00:00 AM	2	
21	10255	RICSU	9	7/12/1996 12:00:00 AM	8/9/1996 12:00:00 AM	7/15/1996 12:00:00 AM	3	
22	10255	RICSU	9	7/12/1996 12:00:00 AM	8/9/1996 12:00:00 AM	7/15/1996 12:00:00 AM	3	
23								

Completed (0.81 s) Columns: 21 Rows: 99+

Query settings: Properties, Name: Orders, Applied steps: Source, Navigation 1, Expanded..., Added cust..., Reordered...

Data destination: No data destination

Step 3- Remove Unnecessary Columns

Unnecessary columns were removed from the Products table to retain only the fields required for further analysis.

Power Query Editor

Home Transform Add column View Help

Close & Apply, Get data, Enter data, Manage parameters, Refresh, Properties, Advanced editor, Export query results, Choose columns, Remove columns, Keep rows, Remove rows, Filter rows, Transform, Combine, Export template

Queries [2]: Orders, Products

Table.TransformColumnTypes(#"Removed other columns", {{{"UnitsInStock", Int64.Type}}})


	ProductID	ProductName	QuantityPerUnit	UnitsInStock
1	1	Chai	10 boxes x 20 bags	39
2	2	Chang	24 - 12 oz bottles	17
3	3	Aniseed Syrup	12 - 550 ml bottles	13
4	4	Chef Anton's Cajun Seasoning	48 - 6 oz jars	53
5	5	Chef Anton's Gumbo Mix	36 boxes	0
6	6	Grandma's Boysenberry Spread	12 - 8 oz jars	120
7	7	Uncle Bob's Organic Dried Pears	12 - 1 lb pkgs.	15
8	8	Northwoods Cranberry Sauce	12 - 12 oz jars	6
9	9	Mishi Kobe Niku	18 - 500 g pkgs.	29
10	10	Ikura	12 - 200 ml jars	31
11	11	Queso Cabrales	1 kg pkg.	22
12	12	Queso Manchego La Pastora	10 - 500 g pkgs.	86
13	13	Konbu	2 kg box	24
14	14	Tofu	40 - 100 g pkgs.	35
15	15	Genen Shouyu	24 - 250 ml bottles	39
16	16	Pavlova	32 - 500 g boxes	29
17	17	Alice Mutton	20 - 1 kg tins	0
18	18	Carnarvon Tigers	16 kg pkg.	42
19	19	Teatime Chocolate Biscuits	10 boxes x 12 pieces	25
20	20	Sir Rodney's Marmalade	30 gift boxes	40
21	21	Sir Rodney's Scones	24 pkgs. x 4 pieces	3
22	22	Gustaf's Knäckebröd	24 - 500 g pkgs.	104
23	23	Trüffelschokolade	12 - 250 g pkgs.	61

Query settings: Properties, Name: Products, Applied steps: Source, Navigation 1, Removed o..., Changed c...

Data destination: No data destination


Step 4 - Verify the Data Type of UnitsInStock

The data type of the UnitsInStock column was verified as Whole Number, which is appropriate for representing stock quantities.




Step 5- Calculate Line Total

A new LineTotal column was created by multiplying UnitPrice by Quantity for each order detail.




Step 6- Rename and Reorder Columns

The columns were verified and arranged in a suitable order for further analysis.




Step 7- Load the Transformed Data

The transformed data was loaded into Power BI Desktop using the Close & Apply option.




Step 8 Create Relationships Between Tables

Relationships were created between the related tables using their common key fields to establish the data model.



Step 9 Result

The ETL process was successfully performed in Power BI by extracting data from an Excel workbook, transforming the data using Power Query, creating the Line Total column, loading the transformed data, and establishing relationships between the tables.



Name : Rushikesh Anil Khairnar

Student ID: 5115301

Roll NO: 21

Subject: Data Engineering

Practical 3: MongoDB Operations

Aim:

To perform single and multiple document insertion and retrieval using MongoDB shell (mongosh).

Step 1: Insert Single Document

Query:

```
db.items.insertOne({ name: "laptop", price: 999 });
```

Output:

```
>_MONGOSH
> db.items.insertOne({ name: "laptop", price: 999 });
< {
  acknowledged: true,
  insertedId: ObjectId('6ac649038fa08d8f4561c384')
}
test>
```

Step 2: View Inserted Document

Query:

```
db.items.find();
```

Output:

```
> db.items.find();
< {
  _id: ObjectId('6ac51e1cd4bbf249c034346a'),
  name: 'laptop',
  price: 999
}
{
  _id: ObjectId('6ac649038fa08d8f4561c384'),
  name: 'laptop',
  price: 999
}
test> |
```

Step 3: Insert Multiple Documents

Query:

```
db.products.insertMany([
  { name: "phone", price: 500, stock: 10 },
  { name: "tablet", price: 300, stock: 5 },
  { name: "watch", price: 150, stock: 0 }
]);
```

Output:

```
> db.products.insertMany([
  { name: "phone", price: 500, stock: 10 },
  { name: "tablet", price: 300, stock: 5 },
  { name: "watch", price: 150, stock: 0 }
]);
< {
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('6ac6498b8fa08d8f4561c385'),
    '1': ObjectId('6ac6498b8fa08d8f4561c386'),
    '2': ObjectId('6ac6498b8fa08d8f4561c387')
  }
}
test>
```

Step 4: View Products in Formatted JSON

Query:

```
db.products.find().pretty();
```

Output:

```
> db.products.find().pretty();
< {
  _id: ObjectId('6ac51e43d4bbf249c034346b'),
  name: 'phone',
  price: 500,
  stock: 10
}
{
  _id: ObjectId('6ac51e43d4bbf249c034346c'),
  name: 'tablet',
  price: 300,
  stock: 5
}
{
  _id: ObjectId('6ac51e43d4bbf249c034346d'),
  name: 'watch',
  price: 150,
  stock: 0
}
{
  _id: ObjectId('6ac6498b8fa08d8f4561c385'),
  name: 'phone',
  price: 500,
  stock: 10
}
```

```
{
  _id: ObjectId('6ac6498b8fa08d8f4561c386'),
  name: 'tablet',
  price: 300,
  stock: 5
}
{
  _id: ObjectId('6ac6498b8fa08d8f4561c387'),
  name: 'watch',
  price: 150,
  stock: 0
}
test>
```

Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 4

Python program demonstrating Noise Elimination, Feature Selection, and Exploratory Data Analysis (EDA) using Pandas, NumPy, Matplotlib, and Scikit-learn.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.feature_selection import VarianceThreshold

# -----
# Sample Dataset
# -----
data = {
    'Age': [20, 21, 22, 23, 24, 100],    # 100 is an outlier (noise)
    'Marks': [75, 80, 85, 90, 95, 20],
    'Attendance': [85, 88, 90, 92, 94, 50],
    'Constant': [1, 1, 1, 1, 1, 1]      # Constant feature
}

df = pd.DataFrame(data)

print("Original Dataset:")
print(df)

# =====
# 1. Noise Elimination (Outlier Removal)
# =====
Q1 = df['Age'].quantile(0.25)
Q3 = df['Age'].quantile(0.75)
IQR = Q3 - Q1

lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR

df_clean = df[(df['Age'] ≥ lower) & (df['Age'] ≤ upper)]

print("\nDataset after Noise Elimination:")
print(df_clean)

# =====
# 2. Feature Selection
# Remove constant or low-variance features
# =====
selector = VarianceThreshold(threshold=0.0)
```

```

selected = selector.fit_transform(df_clean)

selected_columns = df_clean.columns[selector.get_support()]

print("\nSelected Features:")
print(selected_columns)

# =====
# 3. Exploratory Data Analysis (EDA)
# =====

print("\nDataset Information:")
print(df_clean.info())

print("\nStatistical Summary:")
print(df_clean.describe())

print("\nCorrelation Matrix:")
print(df_clean.corr())

# Histogram
df_clean.hist(figsize=(8,6))
plt.suptitle("Histogram of Features")
plt.show()

# Boxplot
df_clean.boxplot(figsize=(8,5))
plt.title("Boxplot of Dataset")
plt.show()

```

Original Dataset:

	Age	Marks	Attendance	Constant
0	20	75	85	1
1	21	80	88	1
2	22	85	90	1
3	23	90	92	1
4	24	95	94	1
5	100	20	50	1

Dataset after Noise Elimination:

	Age	Marks	Attendance	Constant
0	20	75	85	1
1	21	80	88	1
2	22	85	90	1
3	23	90	92	1
4	24	95	94	1

Selected Features:

Index(['Age', 'Marks', 'Attendance'], dtype='object')

Dataset Information:

```

<class 'pandas.core.frame.DataFrame'>
Index: 5 entries, 0 to 4
Data columns (total 4 columns):

```

#	Column	Non-Null Count	Dtype
0	Age	5 non-null	int64
1	Marks	5 non-null	int64
2	Attendance	5 non-null	int64
3	Constant	5 non-null	int64

dtypes: int64(4)

memory usage: 200.0 bytes

None





Statistical Summary:


	Age	Marks	Attendance	Constant
count	5.000000	5.000000	5.000000	5.0
mean	22.000000	85.000000	89.800000	1.0
std	1.581139	7.905694	3.49285	0.0
min	20.000000	75.000000	85.000000	1.0
25%	21.000000	80.000000	88.000000	1.0
50%	22.000000	85.000000	90.000000	1.0
75%	23.000000	90.000000	92.000000	1.0
max	24.000000	95.000000	94.000000	1.0

Correlation Matrix:

	Age	Marks	Attendance	Constant
Age	1.000000	1.000000	0.995893	NaN
Marks	1.000000	1.000000	0.995893	NaN
Attendance	0.995893	0.995893	1.000000	NaN
Constant	NaN	NaN	NaN	NaN

Histogram of Features





Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 5

Extracting Data from APIs and Flat Files Using Python

1. Objective Write a Python script to fetch user data from a public REST API, clean it, and merge it with a flat CSV file containing location details.

2. Code Implementation Create a script named **data_extraction**.

```
import pandas as pd
import requests

def extract_api_data(url):
    """Fetches data from a REST API and returns a DataFrame."""
    try:
        response = requests.get(url, timeout=10)
        response.raise_for_status()
        data = response.json()
        # Normalizing JSON data into a flat table
        return pd.json_normalize(data)
    except requests.exceptions.RequestException as e:
        print(f"API Error: {e}")
        return pd.DataFrame()

def extract_csv_data(file_path):
    """Reads a flat CSV file."""
    try:
        return pd.read_csv(file_path)
    except FileNotFoundError as e:
        print(f"File Error: {e}")
        return pd.DataFrame()

if __name__ == "__main__":
    # 1. API Extraction
    # Using jsonplaceholder to get mock users API data
    api_url = "https://jsonplaceholder.typicode.com/users"
    api_df = extract_api_data(api_url)

    if not api_df.empty:
        api_df = api_df[["id", "name", "email", "company.name"]]
        api_df.rename(columns={"company.name": "company"}, inplace=True)

    # 2. Flat File Extraction (Mock CSV Data)
    csv_file = "locations.csv"
```

```

# Fixed the missing list of IDs to resolve the SyntaxError
pd.DataFrame(
    {
        "id": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
        "city": [
            "New York",
            "London",
            "Paris",
            "Tokyo",
            "Berlin",
            "Delhi",
            "Sydney",
            "Moscow",
            "Cairo",
            "Beijing",
        ],
    },
).to_csv(csv_file, index=False)

csv_df = extract_csv_data(csv_file)

# 3. Data Transformation & Merging
if not api_df.empty and not csv_df.empty:
    merged_df = pd.merge(api_df, csv_df, on="id", how="inner")
    print("\n--- Merged ETL Pipeline Data View ---")
    print(merged_df.head())

# Save to target Data Lake / Warehouse stage
merged_df.to_csv("cleaned_warehouse_profiles.csv", index=False)
print("\nData successfully saved to 'cleaned_warehouse_profiles.csv'")

```

--- Merged ETL Pipeline Data View ---

	id	name	email	company \
0	1	Leanne Graham	Sincere@april.biz	Romaguera-Crona
1	2	Ervin Howell	Shanna@melissa.tv	Deckow-Crist
2	3	Clementine Bauch	Nathan@yesenia.net	Romaguera-Jacobson
3	4	Patricia Lebsack	Julianne.OConner@kory.org	Robel-Corkery
4	5	Chelsey Dietrich	Lucio_Hettinger@annie.ca	Keebler LLC

	city
0	New York
1	London
2	Paris
3	Tokyo
4	Berlin

Data successfully saved to 'cleaned_warehouse_profiles.csv'

Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 6

Setting up Apache Airflow and Creating a DAG
1. Objective Configure a Directed Acyclic Graph (DAG) in Apache Airflow to run daily, executing sequential dependency tasks: a start notifier, an extraction script, and a pipeline success logger.
2. DAG

Implementation Place this code in your Airflow DAGs directory (usually `~/airflow/dags/data_pipeline_dag.py`):
`pythonfrom datetime import datetime, timedelta`

```
!pip install apache-airflow
```

```
Requirement already satisfied: apache-airflow in /usr/local/lib/python3.13/dist-packages (3.3.2)
Requirement already satisfied: apache-airflow-core==3.3.2 in /usr/local/lib/python3.13/dist-packages (from apache-airflow) (3.3.2)
Requirement already satisfied: apache-airflow-task-sdk==1.3.2 in /usr/local/lib/python3.13/dist-packages (from apache-airflow) (1.3.2)
Requirement already satisfied: a2wsgi≥1.10.8 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.10.10)
Requirement already satisfied: aiosqlite≠0.22.0, ≥0.20.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (0.22.1)
Requirement already satisfied: alembic<2.0, ≥1.13.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.20.0)
Requirement already satisfied: apache-airflow-providers-common-compat≥1.7.4 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.20.0)
Requirement already satisfied: apache-airflow-providers-common-io≥1.6.3 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.10.0)
Requirement already satisfied: apache-airflow-providers-common-sql≥1.28.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (2.2.0)
Requirement already satisfied: apache-airflow-providers-smtp≥2.3.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (3.1.0)
Requirement already satisfied: apache-airflow-providers-standard≥1.9.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.20.0)
Requirement already satisfied: argcomplete≥1.10 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (3.7.2)
Requirement already satisfied: asgiref≥2.3.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (3.12.1)
Requirement already satisfied: attrs≠25.2.0, ≥22.1.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (26.1.0)
Requirement already satisfied: cachetools≥6.0.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (6.2.6)
Requirement already satisfied: cadwyn≥6.1.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (7.0.0)
Requirement already satisfied: colorlog≥6.8.2 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (6.12.0)
Requirement already satisfied: cron-descriptor≥1.2.24 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (2.1.1)
Requirement already satisfied: croniter≥6.2.2 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (6.2.4)
Requirement already satisfied: cryptography≥44.0.3 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (50.0.2)
Requirement already satisfied: deprecated≥1.2.13 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (3.0.0)
Requirement already satisfied: dill≥0.2.2 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (0.4.1)
```

Requirement already satisfied: fastapi<0.137.0, ≥0.129.0 in /usr/local/lib/python3.13/dist-packages (from fastapi[standard-no-fastapi-cloud-cli]<0.137.0, ≥0.129.0→apache-airflow-core==3.3.2→apache-airflow) (0.136.3)

Requirement already satisfied: httpx≥0.25.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (0.28.1)

Requirement already satisfied: importlib-metadata≥7.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (9.0.1)

Requirement already satisfied: isoduration≥20.11.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (20.11.0)

Requirement already satisfied: itsdangerous≥2.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (2.2.0)

Requirement already satisfied: jinja2≥3.1.5 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (3.1.6)

Requirement already satisfied: jsonschema≥4.19.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (4.26.0)

Requirement already satisfied: lazy-object-proxy≥1.2.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.12.0)

Requirement already satisfied: libcst≥1.8.2 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.9.0)

Requirement already satisfied: linkify-it-py≥2.0.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (2.2.0)

Requirement already satisfied: lockfile≥0.12.2 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (0.12.2)

Requirement already satisfied: methodtools≥0.4.7 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (0.4.7)

Requirement already satisfied: msgspec≥0.19.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (0.22.0)

Requirement already satisfied: natsort≥8.4.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (8.4.0)

Requirement already satisfied: opentelemetry-api≥1.27.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.45.0)

Requirement already satisfied: opentelemetry-exporter-otlp≥1.27.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.45.0)

Requirement already satisfied: opentelemetry-proto<9999, ≥1.27.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.45.0)

Requirement already satisfied: packaging≥25.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (26.3)

Requirement already satisfied: pathspec≥0.9.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.1.1)

Requirement already satisfied: pendulum≥3.1.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (3.2.0)

Requirement already satisfied: pluggy≥1.5.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (1.6.0)

Requirement already satisfied: psutil≥5.8.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (7.2.2)

Requirement already satisfied: pydantic≥2.11.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (2.13.5)

Requirement already satisfied: pygments≠2.19.0, ≥2.0.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2→apache-airflow) (2.21.0)

Requirement already satisfied: pygtrie $\geq$ 2.5.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2.6.2)

Requirement already satisfied: pyjwt $\geq$ 2.11.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2.15.1)

Requirement already satisfied: python-daemon $\geq$ 3.0.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (3.1.2)

Requirement already satisfied: python-dateutil $\geq$ 2.7.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2.9.0.post0)

Requirement already satisfied: python-slugify $\geq$ 5.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (9.1.2)

Requirement already satisfied: pyyaml $\geq$ 6.0.3 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (6.0.3)

Requirement already satisfied: requests <3 , $\geq$ 2.32.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2.32.4)

Requirement already satisfied: rich-argparse $\geq$ 1.0.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (1.8.0)

Requirement already satisfied: rich $\geq$ 13.6.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (13.9.4)

Requirement already satisfied: setproctitle $\geq$ 1.3.3 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (1.3.8)

Requirement already satisfied: sqlalchemy $\geq$ 2.0.50 in /usr/local/lib/python3.13/dist-packages (from sqlalchemy[asyncio] $\geq$ 2.0.50 $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2.1.1)

Requirement already satisfied: starlette $\geq$ 1.0.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (1.7.0)

Requirement already satisfied: structlog $\geq$ 26.1.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (26.1.0)

Requirement already satisfied: svcs $\geq$ 25.1.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (26.2.0)

Requirement already satisfied: tabulate $\geq$ 0.9.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.9.0)

Requirement already satisfied: tenacity $\geq$ 8.3.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (9.1.4)

Requirement already satisfied: termcolor $\geq$ 3.0.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (3.3.0)

Requirement already satisfied: typing-extensions $\geq$ 4.14.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (4.16.0)

Requirement already satisfied: universal-pathlib $\geq$ 0.3.8 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.3.10)

Requirement already satisfied: uuid6 $\geq$ 2024.7.10 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2025.0.1)

Requirement already satisfied: uvicorn $\geq$ 0.37.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.54.0)

Requirement already satisfied: babel $\geq$ 2.17.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-task-sdk==1.3.2 $\rightarrow$ apache-airflow) (2.18.0)

Requirement already satisfied: fsspec $\geq$ 2023.10.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-task-sdk==1.3.2 $\rightarrow$ apache-airflow) (2025.12.0)

Requirement already satisfied: greenback $\geq$ 1.2.1 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-task-sdk==1.3.2 $\rightarrow$ apache-airflow) (1.3.0)

Requirement already satisfied: Mako in /usr/lib/python3/dist-packages (from alembic <2.0 , $\geq$ 1.13.1 $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (1.3.2.dev0)

Requirement already satisfied: sqlparse $\geq$ 0.5.1 in /usr/local/lib/python3.13/dist-

packages (from apache-airflow-providers-common-sql $\geq$ 1.28.1 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.6.0)
 Requirement already satisfied: more-itertools $\geq$ 9.0.0 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-providers-common-sql $\geq$ 1.28.1 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (10.8.0)
 Requirement already satisfied: aiosmtp $\geq$ 0.1.6 in /usr/local/lib/python3.13/dist-packages (from apache-airflow-providers-smtp $\geq$ 2.3.1 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (5.1.3)
 Requirement already satisfied: typing-inspection $\geq$ 0.4.0 in /usr/local/lib/python3.13/dist-packages (from cadwyn $\geq$ 6.1.1 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.4.4)
 Requirement already satisfied: cffi $\geq$ 2.0.0 in /usr/local/lib/python3.13/dist-packages (from cryptography $\geq$ 44.0.3 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2.1.1)
 Requirement already satisfied: wrapt <3 , $\geq$ 1.16 in /usr/local/lib/python3.13/dist-packages (from deprecated $\geq$ 1.2.13 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2.5.0)
 Requirement already satisfied: annotated-doc $\geq$ 0.0.2 in /usr/local/lib/python3.13/dist-packages (from fastapi $<0.137.0$, $\geq$ 0.129.0 $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq$ 0.129.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.0.5)
 Requirement already satisfied: fastapi-cli $\geq$ 0.0.8 in /usr/local/lib/python3.13/dist-packages (from fastapi-cli[standard-no-fastapi-cloud-cli] $\geq$ 0.0.8; extra = "standard-no-fastapi-cloud-cli" $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq$ 0.129.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.0.32)
 Requirement already satisfied: python-multipart $\geq$ 0.0.18 in /usr/local/lib/python3.13/dist-packages (from fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq$ 0.129.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.0.32)
 Requirement already satisfied: email-validator $\geq$ 2.0.0 in /usr/local/lib/python3.13/dist-packages (from fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq$ 0.129.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2.3.0)
 Requirement already satisfied: pydantic-settings $\geq$ 2.0.0 in /usr/local/lib/python3.13/dist-packages (from fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq$ 0.129.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2.15.0)
 Requirement already satisfied: pydantic-extra-types $\geq$ 2.0.0 in /usr/local/lib/python3.13/dist-packages (from fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq$ 0.129.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2.11.1)
 Requirement already satisfied: greenlet $\neq$ 0.4.17 in /usr/local/lib/python3.13/dist-packages (from greenback $\geq$ 1.2.1 $\rightarrow$ apache-airflow-task-sdk=1.3.2 $\rightarrow$ apache-airflow) (3.5.6)
 Requirement already satisfied: sniffio in /usr/local/lib/python3.13/dist-packages (from greenback $\geq$ 1.2.1 $\rightarrow$ apache-airflow-task-sdk=1.3.2 $\rightarrow$ apache-airflow) (1.3.1)
 Requirement already satisfied: outcome in /usr/local/lib/python3.13/dist-packages (from greenback $\geq$ 1.2.1 $\rightarrow$ apache-airflow-task-sdk=1.3.2 $\rightarrow$ apache-airflow) (1.3.0.post0)
 Requirement already satisfied: anyio in /usr/local/lib/python3.13/dist-packages (from httpx $\geq$ 0.25.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (4.15.1)
 Requirement already satisfied: certifi in /usr/local/lib/python3.13/dist-packages (from httpx $\geq$ 0.25.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2026.7.22)
 Requirement already satisfied: httpcore=1.* in /usr/local/lib/python3.13/dist-packages (from httpx $\geq$ 0.25.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.0.9)
 Requirement already satisfied: idna in /usr/local/lib/python3.13/dist-packages (from httpx $\geq$ 0.25.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (3.20)
 Requirement already satisfied: h11 $\geq$ 0.16 in /usr/local/lib/python3.13/dist-packages (from httpcore=1.* $\rightarrow$ httpx $\geq$ 0.25.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.16.0)
 Requirement already satisfied: zipp $\geq$ 3.20 in /usr/local/lib/python3.13/dist-packages (from importlib-metadata $\geq$ 7.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (4.1.0)
 Requirement already satisfied: arrow $\geq$ 0.15.0 in /usr/local/lib/python3.13/dist-

packages (from isoduration $\geq$ 20.11.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.4.0)

Requirement already satisfied: MarkupSafe $\geq$ 2.0 in /usr/local/lib/python3.13/dist-packages (from jinja2 $\geq$ 3.1.5 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (3.0.3)

Requirement already satisfied: jsonschema-specifications $\geq$ 2023.03.6 in /usr/local/lib/python3.13/dist-packages (from jsonschema $\geq$ 4.19.1 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2025.9.1)

Requirement already satisfied: referencing $\geq$ 0.28.4 in /usr/local/lib/python3.13/dist-packages (from jsonschema $\geq$ 4.19.1 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.37.0)

Requirement already satisfied: rpds-py $\geq$ 0.25.0 in /usr/local/lib/python3.13/dist-packages (from jsonschema $\geq$ 4.19.1 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2026.6.3)

Requirement already satisfied: pyyaml-ft $\geq$ 8.0.0 in /usr/local/lib/python3.13/dist-packages (from libcbst $\geq$ 1.8.2 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (8.0.0)

Requirement already satisfied: wire rope $\geq$ 0.4.7 in /usr/local/lib/python3.13/dist-packages (from methodtools $\geq$ 0.4.7 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.0.0)

Requirement already satisfied: opentelemetry-exporter-otlp-proto-grpc=1.45.0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.45.0)

Requirement already satisfied: opentelemetry-exporter-otlp-proto-http=1.45.0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.45.0)

Requirement already satisfied: googleapis-common-protos $\sim$ 1.57 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-otlp-proto-grpc=1.45.0 $\rightarrow$ opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.75.5)

Requirement already satisfied: grpcio $<$ 2.0.0, $\geq$ 1.66.2 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-otlp-proto-grpc=1.45.0 $\rightarrow$ opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.84.0)

Requirement already satisfied: opentelemetry-exporter-otlp-common=0.66b0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-otlp-proto-grpc=1.45.0 $\rightarrow$ opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.66b0)

Requirement already satisfied: opentelemetry-exporter-otlp-proto-common=1.45.0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-otlp-proto-grpc=1.45.0 $\rightarrow$ opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.45.0)

Requirement already satisfied: opentelemetry-sdk $\sim$ 1.45.0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-otlp-proto-grpc=1.45.0 $\rightarrow$ opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (1.45.0)

Requirement already satisfied: protobuf $<$ 8.0, $\geq$ 5.0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-proto $<$ 9999, $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (6.33.6)

Requirement already satisfied: opentelemetry-exporter-http-transport=0.66b0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-http-transport[urllib3]=0.66b0 $\rightarrow$ opentelemetry-exporter-otlp-proto-http=1.45.0 $\rightarrow$ opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (0.66b0)

Requirement already satisfied: urllib3 $\geq$ 1.26 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-exporter-http-transport[urllib3]=0.66b0 $\rightarrow$ opentelemetry-exporter-otlp-proto-http=1.45.0 $\rightarrow$ opentelemetry-exporter-otlp $\geq$ 1.27.0 $\rightarrow$ apache-airflow-core=3.3.2 $\rightarrow$ apache-airflow) (2.5.0)

Requirement already satisfied: tzdata $\geq$ 2020.1 in /usr/local/lib/python3.13/dist-packages (from pendulum $\geq$ 3.1.0 $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2026.4)

Requirement already satisfied: annotated-types $\geq$ 0.6.0 in /usr/local/lib/python3.13/dist-packages (from pydantic $\geq$ 2.11.0 $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.8.0)

Requirement already satisfied: pydantic-core==2.46.5 in /usr/local/lib/python3.13/dist-packages (from pydantic $\geq$ 2.11.0 $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2.46.5)

Requirement already satisfied: six $\geq$ 1.5 in /usr/local/lib/python3.13/dist-packages (from python-dateutil $\geq$ 2.7.0 $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (1.17.0)

Requirement already satisfied: text-unidecode $\geq$ 1.3 in /usr/local/lib/python3.13/dist-packages (from python-slugify $\geq$ 5.0 $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (1.3)

Requirement already satisfied: charset-normalizer <4 , ≥ 2 in /usr/local/lib/python3.13/dist-packages (from requests <3 , $\geq 2.32.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (3.4.9)

Requirement already satisfied: markdown-it-py $\geq$ 2.2.0 in /usr/local/lib/python3.13/dist-packages (from rich $\geq 13.6.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (4.2.0)

Requirement already satisfied: pathlib-abc $<0.6.0$, $\geq 0.5.1$ in /usr/local/lib/python3.13/dist-packages (from universal-pathlib $\geq 0.3.8$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.5.2)

Requirement already satisfied: click ≥ 7.0 in /usr/local/lib/python3.13/dist-packages (from uvicorn $\geq 0.37.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (8.5.0)

Requirement already satisfied: pycparser in /usr/local/lib/python3.13/dist-packages (from cffi $\geq 2.0.0$ $\rightarrow$ cryptography $\geq 44.0.3$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (3.0)

Requirement already satisfied: dnspython $\geq 2.0.0$ in /usr/local/lib/python3.13/dist-packages (from email-validator $\geq 2.0.0$ $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq 0.129.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (2.8.0)

Requirement already satisfied: typer $\geq 0.16.0$ in /usr/local/lib/python3.13/dist-packages (from fastapi-cli $\geq 0.0.8$ $\rightarrow$ fastapi-cli[standard-no-fastapi-cloud-cli] $\geq 0.0.8$; extra = "standard-no-fastapi-cloud-cli" $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq 0.129.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.27.2)

Requirement already satisfied: rich-toolkit $\geq 0.14.8$ in /usr/local/lib/python3.13/dist-packages (from fastapi-cli $\geq 0.0.8$ $\rightarrow$ fastapi-cli[standard-no-fastapi-cloud-cli] $\geq 0.0.8$; extra = "standard-no-fastapi-cloud-cli" $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq 0.129.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.20.5)

Requirement already satisfied: mdurl ≤ 0.1 in /usr/local/lib/python3.13/dist-packages (from markdown-it-py $\geq 2.2.0$ $\rightarrow$ rich $\geq 13.6.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.1.2)

Requirement already satisfied: python-dotenv $\geq 0.21.0$ in /usr/local/lib/python3.13/dist-packages (from pydantic-settings $\geq 2.0.0$ $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq 0.129.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (1.2.3)

Requirement already satisfied: httptools $\geq 0.8.0$ in /usr/local/lib/python3.13/dist-packages (from uvicorn[standard] $\geq 0.12.0$; extra = "standard-no-fastapi-cloud-cli" $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq 0.129.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.8.0)

Requirement already satisfied: uvloop $\geq 0.15.1$ in /usr/local/lib/python3.13/dist-packages (from uvicorn[standard] $\geq 0.12.0$; extra = "standard-no-fastapi-cloud-cli" $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq 0.129.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.22.1)

Requirement already satisfied: watchfiles ≥ 0.20 in /usr/local/lib/python3.13/dist-packages (from uvicorn[standard] $\geq 0.12.0$; extra = "standard-no-fastapi-cloud-cli" $\rightarrow$ fastapi[standard-no-fastapi-cloud-cli] $<0.137.0$, $\geq 0.129.0$ $\rightarrow$ apache-airflow-core==3.3.2 $\rightarrow$ apache-airflow) (0.22.1)

→apache-airflow) (1.3.0)

Requirement already satisfied: websockets $\geq$ 13.0 in /usr/local/lib/python3.13/dist-packages (from uvicorn[standard] $\geq$ 0.12.0; extra = "standard-no-fastapi-cloud-cli"

→fastapi[standard-no-fastapi-cloud-cli]<0.137.0, $\geq$ 0.129.0→apache-airflow-core==3.3.2

→apache-airflow) (15.0.1)

Requirement already satisfied: opentelemetry-semantic-conventions==0.66b0 in /usr/local/lib/python3.13/dist-packages (from opentelemetry-sdk~=1.45.0

→opentelemetry-exporter-otlp-proto-grpc==1.45.0→opentelemetry-exporter-otlp $\geq$ 1.27.0

→apache-airflow-core==3.3.2→apache-airflow) (0.66b0)

Requirement already satisfied: shellingham $\geq$ 1.3.0 in /usr/local/lib/python3.13/dist-packages (from typer $\geq$ 0.16.0→fastapi-cli $\geq$ 0.0.8→fastapi-cli[standard-no-fastapi-

cloud-cli] $\geq$ 0.0.8; extra = "standard-no-fastapi-cloud-cli"→fastapi[standard-no-fastapi-cloud-cli]<0.137.0, $\geq$ 0.129.0→apache-airflow-core==3.3.2→apache-airflow)

(1.5.4)

```

from datetime import datetime, timedelta
from airflow import DAG
from airflow.operators.bash import BashOperator
from airflow.operators.empty import EmptyOperator

# Default arguments for the workflow pipeline execution
default_args = {
    "owner": "data_engineering_lab",
    "depends_on_past": False,
    "start_date": datetime(2026, 1, 1),
    "email_on_failure": False,
    "email_on_retry": False,
    "retries": 1,
    "retry_delay": timedelta(minutes=5),
}

# Define the DAG context
with DAG(
    "university_etl_orchestration",
    default_args=default_args,
    description="Lab assignment for API and Flat File ETL orchestration",
    schedule=timedelta(days=1),
    catchup=False,
) as dag:

    # Task 1: Pipeline Init Anchor
    start_pipeline = EmptyOperator(task_id="start_pipeline")

    # Task 2: Run Python Extraction Script
    execute_extraction = BashOperator(
        task_id="run_extraction_script",
        bash_command="python3 /absolute/path/to/data_extraction.py",
    )

    # Task 3: Finalize and Log Metrics
    pipeline_complete = BashOperator(
        task_id="log_pipeline_success",
        bash_command='echo "ETL Execution completed successfully at $(date)"',
    )

```

```

2026-10-06T15:06:34.342037Z [warning ] The `airflow.operators.bash.BashOperator`
attribute is deprecated. Please use
`'airflow.providers.standard.operators.bash.BashOperator'`. [py.warnings]
category=DeprecatedImportWarning filename=/tmp/ipykernel_2273/851226352.py lineno=3
2026-10-06T15:06:34.346060Z [warning ] The `airflow.operators.empty.EmptyOperator`
attribute is deprecated. Please use
`'airflow.providers.standard.operators.empty.EmptyOperator'`. [py.warnings]
category=DeprecatedImportWarning filename=/tmp/ipykernel_2273/851226352.py lineno=4

```

Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 7

1.Design an ETL pipeline to extract data from a CSV file, clean it, transform it, and store it in a database.

```
import csv
import sqlite3
import os

# 1. SETUP: Create sample raw CSV
csv_file = 'raw_data.csv'
db_file = 'etl_task1.db'
with open(csv_file, 'w', newline='') as f:
    writer = csv.writer(f)
    writer.writerow(['id', 'name', 'salary'])
    writer.writerow(['1', 'alice smith', '50000'])
    writer.writerow(['2', '', '40000']) # Missing name (needs cleaning)
    writer.writerow(['3', 'bob jones', '60000'])

print("---- Task 1: CSV ETL Pipeline ----")

# 2. EXTRACT
raw_data = []
with open(csv_file, 'r') as f:
    reader = csv.DictReader(f)
    raw_data = list(reader)

# 3. CLEAN & TRANSFORM
clean_data = []
for row in raw_data:
    if not row['name'].strip():
        continue # Clean: Skip missing names

    # Transform: Capitalize name and calculate a 10% tax column
    transformed_row = (
        int(row['id']),
        row['name'].title(),
        float(row['salary']),
        float(row['salary']) * 0.10
    )
    clean_data.append(transformed_row)

# 4. LOAD
conn = sqlite3.connect(db_file)
cursor = conn.cursor()
```

```
cursor.execute('''CREATE TABLE IF NOT EXISTS employees
                  (id INTEGER, name TEXT, salary REAL, tax REAL)''')
cursor.executemany('INSERT INTO employees VALUES (?, ?, ?, ?)', clean_data)
conn.commit()

# 5. VISIBLE OUTPUT
cursor.execute('SELECT * FROM employees')
print("Database contents after ETL:")
for row in cursor.fetchall():
    print(row)

# Cleanup
conn.close()
os.remove(csv_file)
os.remove(db_file)
```

```
--- Task 1: CSV ETL Pipeline ---
Database contents after ETL:
(1, 'Alice Smith', 50000.0, 5000.0)
(3, 'Bob Jones', 60000.0, 6000.0)
```

2. Extract data from multiple CSV files and combine them into a single dataset.

```

import csv
import os

# 1. SETUP: Create multiple CSVs
files = ['region_a.csv', 'region_b.csv']
with open(files[0], 'w', newline='') as f:
    csv.writer(f).writerows([[ 'id', 'product'], [ '101', 'Laptop'], [ '102', 'Mouse']])
with open(files[1], 'w', newline='') as f:
    csv.writer(f).writerows([[ 'id', 'product'], [ '103', 'Keyboard'], [ '104',
'Monitor']])

print("\n--- Task 2: Combine Multiple CSVs ---")

# 2. EXTRACT AND COMBINE
combined_dataset = []
header_extracted = False

for file in files:
    with open(file, 'r') as f:
        reader = csv.reader(f)
        header = next(reader)
        if not header_extracted:
            combined_dataset.append(header)
            header_extracted = True

        # Append data rows
        for row in reader:
            combined_dataset.append(row)

# 3. VISIBLE OUTPUT
print("Combined Dataset:")
for row in combined_dataset:
    print(row)

# Cleanup
for file in files:
    os.remove(file)

```

```

--- Task 2: Combine Multiple CSVs ---
Combined Dataset:
['id', 'product']
['101', 'Laptop']
['102', 'Mouse']
['103', 'Keyboard']
['104', 'Monitor']

```

3. Build a pipeline to extract JSON data, transform selected fields, and load it into a relational database.

```

import json
import sqlite3
import os

```

```

# 1. SETUP: Create JSON file
json_file = 'orders.json'
db_file = 'etl_task3.db'
data = [
    {"order_id": "A1", "details": {"customer": "John", "total_cents": 2500}, "date":
"2026-10-06"},
    {"order_id": "A2", "details": {"customer": "Jane", "total_cents": 5000}, "date":
"2026-10-07"}
]
with open(json_file, 'w') as f:
    json.dump(data, f)

print("\n--- Task 3: JSON ETL Pipeline ---")

# 2. EXTRACT
with open(json_file, 'r') as f:
    json_data = json.load(f)

# 3. TRANSFORM
transformed_data = []
for item in json_data:
    # Flatten nested JSON, convert cents to dollars, format date string
    transformed_data.append((
        item['order_id'],
        item['details']['customer'],
        item['details']['total_cents'] / 100.0, # Transform to dollars
        item['date'].replace("-", "/")         # Transform date format
    ))

# 4. LOAD
conn = sqlite3.connect(db_file)
cursor = conn.cursor()
cursor.execute('''CREATE TABLE IF NOT EXISTS orders
                  (order_id TEXT, customer TEXT, total_dollars REAL, order_date
TEXT)''')
cursor.executemany('INSERT INTO orders VALUES (?, ?, ?, ?)', transformed_data)
conn.commit()

# 5. VISIBLE OUTPUT
cursor.execute('SELECT * FROM orders')
print("Database contents after JSON ETL:")
for row in cursor.fetchall():
    print(row)

# Cleanup
conn.close()
os.remove(json_file)
os.remove(db_file)

```

--- Task 3: JSON ETL Pipeline ---
Database contents after JSON ETL:

```
('A1', 'John', 25.0, '2026/10/06')  
( 'A2', 'Jane', 50.0, '2026/10/07')
```

4.Design an ETL process to identify and remove invalid records.

```

print("\n--- Task 4: Remove Invalid Records ---")

# 1. EXTRACT (Simulated raw data)
raw_users = [
    {"user_id": 1, "age": 25, "email": "alice@test.com"},
    {"user_id": 2, "age": -5, "email": "bob@test.com"}, # Invalid age
    {"user_id": 3, "age": 30, "email": "charlie_no_at_symbol.com"}, # Invalid email
    {"user_id": 4, "age": 40, "email": "david@test.com"}
]

valid_records = []
invalid_records = []

# 2. CLEAN (Identify and split)
for user in raw_users:
    is_valid = True
    reason = []

    if user['age'] < 0 or user['age'] > 120:
        is_valid = False
        reason.append("Age out of bounds")

    if "@" not in user['email']:
        is_valid = False
        reason.append("Invalid email format")

    if is_valid:
        valid_records.append(user)
    else:
        user['error_reason'] = " | ".join(reason)
        invalid_records.append(user)

# 3. VISIBLE OUTPUT
print("Valid Records Retained:")
for r in valid_records:
    print(f" {r}")

print("\nInvalid Records Removed:")
for r in invalid_records:
    print(f" {r}")

```

--- Task 4: Remove Invalid Records ---

Valid Records Retained:

```

{'user_id': 1, 'age': 25, 'email': 'alice@test.com'}
{'user_id': 4, 'age': 40, 'email': 'david@test.com'}

```

Invalid Records Removed:

```

{'user_id': 2, 'age': -5, 'email': 'bob@test.com', 'error_reason': 'Age out of
bounds'}
{'user_id': 3, 'age': 30, 'email': 'charlie_no_at_symbol.com', 'error_reason':
'Invalid email format'}

```

5. Create a pipeline that performs data validation before loading data into the target database.

```
import sqlite3
import os

print("\n--- Task 5: Pre-Load Validation ---")

db_file = 'etl_task5.db'
conn = sqlite3.connect(db_file)
cursor = conn.cursor()
cursor.execute('CREATE TABLE products (id INTEGER, name TEXT, price REAL)')

incoming_data = [
    (1, "Table", 150.0),
    (2, "Chair", -20.0), # Fails price validation
    (3, None, 50.0)      # Fails name validation
]

# Validation Function
def validate_product(record):
    product_id, name, price = record
    if not name:
        return False, "Missing name"
    if price < 0:
        return False, "Negative price"
    return True, "Valid"

# Process and Load
successful_inserts = 0
for record in incoming_data:
    is_valid, msg = validate_product(record)
    if is_valid:
        cursor.execute('INSERT INTO products VALUES (?, ?, ?)', record)
        successful_inserts += 1
    else:
        print(f"Validation Failed for record {record}: {msg}")

conn.commit()

# VISIBLE OUTPUT
print(f"\nSuccessfully loaded {successful_inserts} valid records.")
cursor.execute('SELECT * FROM products')
for row in cursor.fetchall():
    print(row)

# Cleanup
conn.close()
os.remove(db_file)
```

--- Task 5: Pre-Load Validation ---

Validation Failed for record (2, 'Chair', -20.0): Negative price

Validation Failed for record (3, None, 50.0): Missing name

Successfully loaded 1 valid records.
(1, 'Table', 150.0)

6. Implement incremental data loading instead of loading the entire dataset every time.

```
import sqlite3
import os

print("\n--- Task 6: Incremental Loading ---")

db_file = 'etl_task6.db'
conn = sqlite3.connect(db_file)
cursor = conn.cursor()

# 1. SETUP: Create table and load initial historical data
cursor.execute('''CREATE TABLE inventory
                  (item_id INTEGER PRIMARY KEY, item_name TEXT, qty INTEGER)''')
cursor.execute("INSERT INTO inventory VALUES (1, 'Pen', 100)")
cursor.execute("INSERT INTO inventory VALUES (2, 'Pencil', 200)")
conn.commit()

print("Initial Database State:")
for row in cursor.execute('SELECT * FROM inventory').fetchall():
    print(f" {row}")

# 2. EXTRACT: New batch of data (Contains 1 update, 1 new item, 1 unchanged)
new_batch = [
    (1, 'Pen', 100),      # Unchanged
    (2, 'Pencil', 250),  # Incremental Update (Quantity changed)
    (3, 'Eraser', 50)    # Incremental Insert (New item)
]

# 3. INCREMENTAL LOAD (UPSERT Strategy)
# INSERT OR REPLACE automatically updates the row if the PRIMARY KEY (item_id)
# already exists.
cursor.executemany('''
    INSERT OR REPLACE INTO inventory (item_id, item_name, qty)
    VALUES (?, ?, ?)
''', new_batch)
conn.commit()

# 4. VISIBLE OUTPUT
print("\nDatabase State After Incremental Load (Notice ID 2 updated, ID 3 added, no
duplicates):")
for row in cursor.execute('SELECT * FROM inventory').fetchall():
    print(f" {row}")

# Cleanup
conn.close()
os.remove(db_file)
```

--- Task 6: Incremental Loading ---

Initial Database State:

(1, 'Pen', 100)
(2, 'Pencil', 200)

Database State After Incremental Load (Notice ID 2 updated, ID 3 added, no duplicates):

(1, 'Pen', 100)
(2, 'Pencil', 250)
(3, 'Eraser', 50)

Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 8

1. Write a PySpark program to read a CSV dataset and display its schema.

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import col, avg, round
import os

# =====
# 0. SETUP: Initialize Spark and Mock Data
# =====
# Initialize a SparkSession
spark = SparkSession.builder \
    .appName("PySpark_Data_Operations") \
    .master("local[*]") \
    .getOrCreate()

# Create mock sales data (contains a duplicate for transaction_id 3)
with open("sales.csv", "w") as f:
    f.write("transaction_id,product_id,quantity,sale_amount\n")
    f.write("1,101,2,50.0\n")
    f.write("2,102,1,120.0\n")
    f.write("3,101,4,100.0\n")
    f.write("3,101,4,100.0\n") # Duplicate row
    f.write("4,103,1,30.0\n")
    f.write("5,101,0,0.0\n") # Invalid quantity to filter out

# Create mock products data
with open("products.csv", "w") as f:
    f.write("product_id,product_name,category\n")
    f.write("101,Wireless Mouse,Electronics\n")
    f.write("102,Mechanical Keyboard,Electronics\n")
    f.write("103,Coffee Mug,Home & Kitchen\n")

# =====
# 1. READ CSV & DISPLAY SCHEMA
# =====
print("\n--- 1. Reading Data and Displaying Schema ---")
# Read CSV files into DataFrames with schema inference
sales_df = spark.read.csv("sales.csv", header=True, inferSchema=True)
products_df = spark.read.csv("products.csv", header=True, inferSchema=True)

# Display the schema of the sales DataFrame
sales_df.printSchema()

```

```

--- 1. Reading Data and Displaying Schema ---
root
 |-- transaction_id: integer (nullable = true)
 |-- product_id: integer (nullable = true)
 |-- quantity: integer (nullable = true)
 |-- sale_amount: double (nullable = true)

```

2. Implement a PySpark program to remove duplicate records.

```
# 3. REMOVE DUPLICATE RECORDS
# =====
print("\n--- 2. Removing Duplicates ---")
print(f"Original row count: {sales_df.count()}")

# Drop duplicate rows across all columns
clean_sales_df = sales_df.dropDuplicates()

print(f"Row count after dropping duplicates: {clean_sales_df.count()}")
clean_sales_df.show()
```

```
--- 2. Removing Duplicates ---
Original row count: 6
Row count after dropping duplicates: 5
+-----+-----+-----+-----+
|transaction_id|product_id|quantity|sale_amount|
+-----+-----+-----+-----+
|          1|      101|        2|       50.0|
|          4|      103|        1|       30.0|
|          2|      102|        1|      120.0|
|          3|      101|        4|      100.0|
|          5|      101|        0|        0.0|
+-----+-----+-----+-----+
```

3. Perform filtering, grouping, and aggregation operations using PySpark DataFrames.

```
# =====
# 2. FILTERING
# =====
print("\n--- 3. Filtering Operations ---")
# Filter out records where quantity is 0 or less
filtered_sales_df = clean_sales_df.filter(col("quantity") > 0)
filtered_sales_df.show()
```

```
--- 3. Filtering Operations ---
+-----+-----+-----+-----+
|transaction_id|product_id|quantity|sale_amount|
+-----+-----+-----+-----+
|          1|      101|        2|       50.0|
|          4|      103|        1|       30.0|
|          2|      102|        1|      120.0|
|          3|      101|        4|      100.0|
+-----+-----+-----+-----+
```

4. Perform a join between two datasets using PySpark.

```
# 4. PERFORM A JOIN BETWEEN TWO DATASETS
# =====
print("\n--- 4. Joining DataFrames ---")
# Perform an inner join on 'product_id'
joined_df = filtered_sales_df.join(products_df, on="product_id", how="inner")
joined_df.show()
```

--- 4. Joining DataFrames ---

product_id	transaction_id	quantity	sale_amount	product_name	category
101	1	2	50.0	Wireless Mouse	Electronics
103	4	1	30.0	Coffee Mug	Home & Kitchen
102	2	1	120.0	Mechanical Keyboard	Electronics
101	3	4	100.0	Wireless Mouse	Electronics

5. Write a PySpark program to calculate average sales by product category.

```
# 5. GROUPING & AGGREGATION (Avg Sales by Category)
# =====
print("\n--- 5. Average Sales by Product Category ---")
# Group by category, calculate the average sale amount, and round to 2 decimal places
avg_sales_df = joined_df.groupBy("category") \
    .agg(
        round(avg("sale_amount"), 2).alias("average_sales")
    )

avg_sales_df.show()

# =====
# CLEANUP
# =====
# Stop the Spark session and remove mock files
spark.stop()
os.remove("sales.csv")
os.remove("products.csv")
```

--- 5. Average Sales by Product Category ---

category	average_sales
Home & Kitchen	30.0
Electronics	90.0

Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 9

1. Build an end-to-end pipeline: CSV → Python/Pandas → Data Cleaning → SQL Database → Report.

```
import pandas as pd
import sqlite3
import os

print("--- Task 1: End-to-End CSV to SQL Pipeline ---")

csv_file = 'raw_sales.csv'
db_file = 'task1_pipeline.db'

# 1. SETUP: Create messy raw data
with open(csv_file, 'w') as f:
    f.write("transaction_id,date,amount,category\n")
    f.write("1,2026-10-01,150.50,Electronics\n")
    f.write("2,,200.00,Clothing\n")          # Missing date
    f.write("3,2026-10-03,ERROR,Electronics\n") # Invalid amount string
    f.write("4,2026-10-04,99.99,Clothing\n")

# 2. EXTRACT: Read with Pandas
df = pd.read_csv(csv_file)

# 3. TRANSFORM & CLEAN
# Drop rows where 'date' is missing
df = df.dropna(subset=['date'])
# Convert 'amount' to numeric, forcing errors to NaN, then drop those NaNs
df['amount'] = pd.to_numeric(df['amount'], errors='coerce')
df = df.dropna(subset=['amount'])
# Ensure proper date formatting
df['date'] = pd.to_datetime(df['date']).dt.strftime('%Y-%m-%d')

# 4. LOAD: Write to SQL Database
conn = sqlite3.connect(db_file)
df.to_sql('cleaned_sales', conn, if_exists='replace', index=False)

# 5. REPORT: Generate summary report using SQL
report_query = """
    SELECT category, COUNT(transaction_id) as total_sales, SUM(amount) as
    total_revenue
    FROM cleaned_sales
    GROUP BY category
    """
```

```
report_df = pd.read_sql(report_query, conn)

print("\nFinal Aggregated Report:")
print(report_df.to_string(index=False))

# Cleanup
conn.close()
os.remove(csv_file)
os.remove(db_file)
```

--- Task 1: End-to-End CSV to SQL Pipeline ---

Final Aggregated Report:

category	total_sales	total_revenue
Clothing	1	99.99
Electronics	1	150.50

2. Build an ETL pipeline for an e-commerce dataset containing customers, orders, products, and payments.

```
import pandas as pd
import sqlite3
import os

print("\n--- Task 2: E-Commerce ETL Pipeline ---")

db_file = 'task2_ecommerce.db'

# 1. EXTRACT: Simulated extraction of 4 different entities
customers = pd.DataFrame({
    'cust_id': [1, 2],
    'name': ['Alice', 'Bob']
})
products = pd.DataFrame({
    'prod_id': [101, 102],
    'price': [1200.0, 50.0] # 101=Laptop, 102=Mouse
})
orders = pd.DataFrame({
    'order_id': [1001, 1002, 1003],
    'cust_id': [1, 1, 2],
    'prod_id': [101, 102, 102],
    'qty': [1, 2, 1]
})
payments = pd.DataFrame({
    'order_id': [1001, 1002, 1003],
    'status': ['Completed', 'Failed', 'Completed']
})

# 2. TRANSFORM
# Join Orders with Payments to filter only successful transactions
merged = orders.merge(payments, on='order_id')
successful_orders = merged[merged['status'] == 'Completed']
```

```

# Join with Products to calculate total transaction value
successful_orders = successful_orders.merge(products, on='prod_id')
successful_orders['total_value'] = successful_orders['qty'] *
successful_orders['price']

# Join with Customers to get names and aggregate revenue per customer
final_merged = successful_orders.merge(customers, on='cust_id')
customer_revenue = final_merged.groupby('name')['total_value'].sum().reset_index()
customer_revenue.rename(columns={'total_value': 'lifetime_revenue'}, inplace=True)

# 3. LOAD
conn = sqlite3.connect(db_file)
customer_revenue.to_sql('customer_revenue_report', conn, if_exists='replace',
index=False)

# 4. VISIBLE OUTPUT
print("\nFinal Customer Revenue Table (Notice Bob's failed order is excluded):")
print(pd.read_sql("SELECT * FROM customer_revenue_report",
conn).to_string(index=False))

# Cleanup
conn.close()
os.remove(db_file)

```

--- Task 2: E-Commerce ETL Pipeline ---

Final Customer Revenue Table (Notice Bob's failed order is excluded):

name	lifetime_revenue
Alice	1200.0
Bob	50.0

3.Design a data pipeline that handles both historical and newly arriving data.

```

import pandas as pd
import sqlite3
import os

print("\n--- Task 3: Historical & Arriving Data Pipeline ---")

db_file = 'task3_incremental.db'
conn = sqlite3.connect(db_file)
cursor = conn.cursor()

# 1. SETUP: Create table and insert HISTORICAL data
cursor.execute('''CREATE TABLE user_profiles
                  (user_id INTEGER PRIMARY KEY, email TEXT, status TEXT)''')
historical_data = [
    (1, 'alice@test.com', 'Active'),
    (2, 'bob@test.com', 'Active')
]
cursor.executemany('INSERT INTO user_profiles VALUES (?, ?, ?)', historical_data)
conn.commit()

```

```

print("Initial Historical Data in DB:")
print(pd.read_sql("SELECT * FROM user_profiles", conn).to_string(index=False))

# 2. EXTRACT: Newly arriving data (Contains an update for ID 2, and a new record for ID 3)
new_data_df = pd.DataFrame({
    'user_id': [2, 3],
    'email': ['bob_updated@test.com', 'charlie@test.com'],
    'status': ['Inactive', 'Active']
})

# 3. TRANSFORM & LOAD: Incremental Upsert Strategy
# Step A: Load new data into a temporary staging table
new_data_df.to_sql('staging_profiles', conn, if_exists='replace', index=False)

# Step B: Perform UPSERT (Insert or Replace) from staging to main table based on Primary Key
cursor.execute('''
    INSERT OR REPLACE INTO user_profiles (user_id, email, status)
    SELECT user_id, email, status FROM staging_profiles
''')
conn.commit()

# Step C: Drop staging table
cursor.execute('DROP TABLE staging_profiles')
conn.commit()

# 4. VISIBLE OUTPUT
print("\nDatabase State After Incremental Load (ID 2 Updated, ID 3 Inserted):")
print(pd.read_sql("SELECT * FROM user_profiles", conn).to_string(index=False))

# Cleanup
conn.close()
os.remove(db_file)

```

--- Task 3: Historical & Arriving Data Pipeline ---

Initial Historical Data in DB:

user_id	email	status
1	alice@test.com	Active
2	bob@test.com	Active

Database State After Incremental Load (ID 2 Updated, ID 3 Inserted):

user_id	email	status
1	alice@test.com	Active
2	bob_updated@test.com	Inactive
3	charlie@test.com	Active

Name: Rushikesh Anil Khairnar

Student: 5115301

Roll NO: 21

College: B. K. Birla College

1. Practical 10: End-to-End Data Engineering Mini Project

This comprehensive, standalone mini-project fulfills the requirements of Practical 10. While originally designed for a group of two, this scope has been expanded and adapted for a single developer to showcase a complete Data Engineering lifecycle.

The project simulates a retail company's pipeline. It includes:

- 1.Data Ingestion:** Reading from diverse sources (CSV and JSON).
- 2.Data Cleaning & Transformation:** Handling missing values, standardizing dates, and calculating derived metrics.
- 3.Data Warehouse Design:** Implementing a Star Schema (Fact and Dimension tables).
- 4.Storage:** Loading the transformed data into a relational SQLite database.
- 5.Reporting:** Querying the warehouse to generate a business report and a data visualization.

2. Complete Project Code

It is completely self-contained; it will generate the necessary dummy raw data, process it through the entire pipeline, build the database, and output the final reports.

```
import pandas as pd
import sqlite3
import json
import os
import matplotlib.pyplot as plt

print("=== STARTING END-TO-END DATA ENGINEERING PIPELINE ===")

# =====
# PHASE 0: SETUP & RAW DATA GENERATION
# =====
print("\n[Phase 0] Generating simulated raw data (CSVs and JSON)...")

# 1. Customers Data (CSV)
customers_data = """customer_id,first_name,last_name,region
1,Alice,Smith,North
2,Bob,Jones,South
3,Charlie,Brown,East
4,Diana,Prince,West
```

```

5,Eve,,North"" # Intentional missing last name
with open("raw_customers.csv", "w") as f:
    f.write(customers_data)

# 2. Products Data (JSON)
products_data = [
    {"product_id": 101, "product_name": "Laptop", "category": "Electronics", "price":
1200.00},
    {"product_id": 102, "product_name": "Desk Chair", "category": "Furniture",
"price": 150.50},
    {"product_id": 103, "product_name": "Mouse", "category": "Electronics", "price":
25.00},
    {"product_id": 104, "product_name": "Notebook", "category": "Stationery",
"price": 5.00}
]
with open("raw_products.json", "w") as f:
    json.dump(products_data, f)

# 3. Transactions Data (CSV) - Contains anomalies (missing dates, negative qty)
transactions_data = ""tx_id,date,customer_id,product_id,quantity
1001,2026-10-01,1,101,1
1002,2026-10-02,2,102,2
1003,,3,103,1
1004,2026-10-02,1,103,2
1005,2026-10-03,4,104,-5
1006,2026-10-04,5,101,1""
with open("raw_transactions.csv", "w") as f:
    f.write(transactions_data)

# =====
# PHASE 1: DATA INGESTION (EXTRACTION)
# =====
print("\n[Phase 1] Extracting data from multiple sources...")
df_customers = pd.read_csv("raw_customers.csv")
df_products = pd.read_json("raw_products.json")
df_transactions = pd.read_csv("raw_transactions.csv")

# =====
# PHASE 2: DATA CLEANING & TRANSFORMATION
# =====
print("\n[Phase 2] Cleaning and Transforming data...")

# Clean Customers: Fill missing last names, create full name
df_customers['last_name'] = df_customers['last_name'].fillna("Unknown")
df_customers['full_name'] = df_customers['first_name'] + " " +
df_customers['last_name']
df_customers = df_customers[['customer_id', 'full_name', 'region']] # Select final
columns

# Clean Transactions: Remove invalid rows
initial_tx_count = len(df_transactions)
df_transactions = df_transactions.dropna(subset=['date']) # Drop missing dates

```

```

df_transactions = df_transactions[df_transactions['quantity'] > 0] # Remove negative quantities
print(f" → Dropped {initial_tx_count - len(df_transactions)} invalid transaction records.")

# Transformation: Derive Date Dimension elements
df_transactions['date'] = pd.to_datetime(df_transactions['date'])
df_date_dim = pd.DataFrame({
    'date_id': df_transactions['date'].dt.strftime('%Y%m%d').astype(int),
    'full_date': df_transactions['date'].dt.date,
    'year': df_transactions['date'].dt.year,
    'month': df_transactions['date'].dt.month,
    'day': df_transactions['date'].dt.day
}).drop_duplicates()

# Transformation: Map date_id back to transactions and calculate Total Sales Amount
df_transactions['date_id'] = df_transactions['date'].dt.strftime('%Y%m%d').astype(int)
df_transactions = df_transactions.merge(df_products[['product_id', 'price']],
on='product_id', how='left')
df_transactions['total_amount'] = df_transactions['quantity'] * df_transactions['price']

# Prepare Final Fact Table
df_fact_sales = df_transactions[['tx_id', 'date_id', 'customer_id', 'product_id', 'quantity', 'total_amount']]

# =====
# PHASE 3: DATA WAREHOUSE DESIGN & STORAGE
# =====
print("\n[Phase 3] Loading Star Schema into SQLite Data Warehouse...")
db_name = "enterprise_data_warehouse.db"
conn = sqlite3.connect(db_name)

# Load Dimensions
df_customers.to_sql("dim_customer", conn, if_exists="replace", index=False)
df_products.to_sql("dim_product", conn, if_exists="replace", index=False)
df_date_dim.to_sql("dim_date", conn, if_exists="replace", index=False)

# Load Fact
df_fact_sales.to_sql("fact_sales", conn, if_exists="replace", index=False)
print(f" → Successfully loaded {len(df_fact_sales)} records into fact_sales.")

# =====
# PHASE 4: REPORTING & VISUALIZATION
# =====
print("\n[Phase 4] Generating Business Reports from Data Warehouse...")

# Report 1: Sales by Category (SQL JOIN across Star Schema)
query_sales_by_category = """
    SELECT
        p.category,

```

```

        SUM(f.quantity) as items_sold,
        SUM(f.total_amount) as total_revenue
    FROM fact_sales f
    JOIN dim_product p ON f.product_id = p.product_id
    GROUP BY p.category
    ORDER BY total_revenue DESC
"""
df_report1 = pd.read_sql(query_sales_by_category, conn)
print("\n--- REPORT 1: Revenue by Category ---")
print(df_report1.to_string(index=False))

# Report 2: Top Customers by Region
query_top_customers = """
    SELECT
        c.full_name,
        c.region,
        SUM(f.total_amount) as lifetime_value
    FROM fact_sales f
    JOIN dim_customer c ON f.customer_id = c.customer_id
    GROUP BY c.full_name, c.region
    ORDER BY lifetime_value DESC
"""
df_report2 = pd.read_sql(query_top_customers, conn)
print("\n--- REPORT 2: Top Customers ---")
print(df_report2.to_string(index=False))

# Visualization: Bar Chart of Sales by Category
plt.figure(figsize=(8, 5))
plt.bar(df_report1['category'], df_report1['total_revenue'], color=['#4C72B0',
'#55A868'])
plt.title('Total Revenue by Product Category')
plt.xlabel('Category')
plt.ylabel('Revenue ($)')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
chart_filename = "revenue_report_chart.png"
plt.savefig(chart_filename)
print(f"\n → Visualization saved locally as '{chart_filename}'")

# =====
# CLEANUP
# =====
conn.close()
os.remove("raw_customers.csv")
os.remove("raw_products.json")
os.remove("raw_transactions.csv")
print("\n≡ PIPELINE EXECUTION COMPLETE ≡")

```

≡ STARTING END-TO-END DATA ENGINEERING PIPELINE ≡

[Phase 0] Generating simulated raw data (CSVs and JSON)...

[Phase 1] Extracting data from multiple sources...

[Phase 2] Cleaning and Transforming data...

→ Dropped 2 invalid transaction records.

[Phase 3] Loading Star Schema into SQLite Data Warehouse...

→ Successfully loaded 4 records into fact_sales.

[Phase 4] Generating Business Reports from Data Warehouse...

--- REPORT 1: Revenue by Category ---


category	items_sold	total_revenue
Electronics	4	2450.0
Furniture	2	301.0

--- REPORT 2: Top Customers ---

full_name	region	lifetime_value
Alice Smith	North	1250.0
Eve Unknown	North	1200.0
Bob Jones	South	301.0

→ Visualization saved locally as 'revenue_report_chart.png'

≡ PIPELINE EXECUTION COMPLETE ≡



```
# Optional: Export Star Schema to CSV for Power BI / Tableau
df_fact_sales.to_csv("fact_sales.csv", index=False)
df_customers.to_csv("dim_customer.csv", index=False)
df_products.to_csv("dim_product.csv", index=False)
df_date_dim.to_csv("dim_date.csv", index=False)
print(" → Exported Star Schema to CSV for BI Tools.")
```

→ Exported Star Schema to CSV for BI Tools.

3. Key Architectural Concepts Addressed

Data Ingestion: Handled mixed formats (CSV & JSON) and standardized them into Pandas DataFrames.

Data Cleaning: Removed corrupted transactions (null dates, negative quantities) and handled missing customer strings.

Data Warehouse Design (Star Schema): Modeled the data into a central fact_sales table surrounded by dim_customer, dim_product, and dim_date tables. This allows for highly optimized analytical querying.

Storage: Utilized an embedded SQLite database (enterprise_data_warehouse.db) to persist the architecture.

4. Power BI Dashboard

